

國立台灣科技大學電子工程系
基於深度學習之影像辨識
期末報告


人臉生成 with DCGAN

姓名:劉家瑜
學號:M11102150

中華民國 111 年 12 月 28 日

一、摘要

GAN(Generative Adversarial Network)在前幾年與近期都是非常熱門的領域，從下圖歷年論文數量可以觀察出來。



參考自 G. Ian et al., "Generative adversarial networks,"

Communications of the ACM, vol. 63, no. 11, pp. 139–144, 2020.,而這邊

我的實作是 DCGAN(Deep Convolutional Generative Adversarial

Networks)，顧名思義該神經網路模型採用的是深度捲積網路，由多

層捲積層堆疊而成，它的好處在於可以提升 GAN 訓練的穩定性及

生成結果的質量。

二、簡介

生成對抗網路的概念非常簡單，就像一個遊戲有兩個角色，一個是偽造者(counterfeiter)，他負責製造假鈔，而另一個角色則是警察，需要不斷從偽造者那邊拿到假鈔並判斷究竟是真是假，然後偽造者會根據警察判斷的結果回饋不斷改良假鈔的品質，最後假鈔的品質被改良到和真鈔很相似，因而很難辨別。

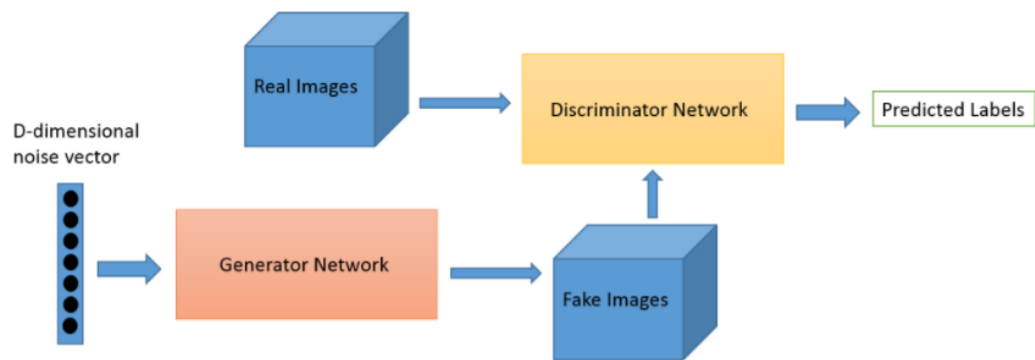
硬體規格:

CPU:i7-12700

GPU:RTX 3090

三、設計原理

GAN 分成兩個網路，Generator 跟 Discriminator，而我的期末報告是做人臉生成，因此我的 Generator 要不斷製造出假的人臉，而 Discriminator 則要負責去判斷到底是 fake image 還是 real image，而 Generator 會根據反饋不斷修正並製造 fake image，最後得出來的人臉就像是真的一樣，但其實那個人並不真實存在，一切都只是由 random noise 經過 Generator 製造出來的虛幻人物而已。



Generator model:

```

Generator(
  (main): Sequential(
    (0): ConvTranspose2d(100, 512, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (5): ReLU(inplace=True)
    (6): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (7): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (8): ReLU(inplace=True)
    (9): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (10): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (11): ReLU(inplace=True)
    (12): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (13): Tanh()
  )
)
  
```

Discriminator model:

```

Discriminator(
  (main): Sequential(
    (0): Conv2d(3, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): LeakyReLU(negative_slope=0.2, inplace=True)
    (2): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (4): LeakyReLU(negative_slope=0.2, inplace=True)
    (5): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (6): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (7): LeakyReLU(negative_slope=0.2, inplace=True)
    (8): Conv2d(256, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (9): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (10): LeakyReLU(negative_slope=0.2, inplace=True)
    (11): Conv2d(512, 1, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (12): Sigmoid()
  )
)

```

Dataset:

使用網路上的 Large-scale CelebFaces Attributes (CelebA)

Dataset，這個資料集共有 202,599 張面部圖像，收錄了 10,177

位知名人物，並且該數據集中涵蓋了巨大的姿態變化和背景雜

訊，這個資料集除了拿來訓練 GAN 模型之外，也可以拿來做為

臉部辨識或是臉部偵測的訓練或測試資料集，下圖皆為範例圖

片。




四、實驗成果

我做了很多組實驗，分別使用不同的 epochs 與不同的 batch size 和 learning rate 並分別把它們記錄下來。

Batch size = 128, Epochs = 5, Learning Rate = 0.0002


```
... Starting Training Loop...  
[0/5][0/28]   Loss_D: 2.0722 Loss_G: 5.2568 D(x): 0.4722 D(G(z)): 0.6326 / 0.0078  
[1/5][0/28]   Loss_D: 0.2588 Loss_G: 12.6711 D(x): 0.9797 D(G(z)): 0.1855 / 0.0000  
[2/5][0/28]   Loss_D: 0.4460 Loss_G: 21.8817 D(x): 0.7986 D(G(z)): 0.0000 / 0.0000  
[3/5][0/28]   Loss_D: 0.5167 Loss_G: 29.6986 D(x): 0.8039 D(G(z)): 0.0000 / 0.0000  
[4/5][0/28]   Loss_D: 1.2115 Loss_G: 11.2988 D(x): 0.9599 D(G(z)): 0.5839 / 0.0001
```



可以看出經過 5 個 epochs 後，從原本的 random noise 已經漸漸有一點人臉的輪廓。

Batch size = 128, Epochs = 10, Learning Rate = 0.0002

```
... Starting Training Loop...
[0/10][0/28]   Loss_D: 2.0722  Loss_G: 5.2567  D(x): 0.4722  D(G(z)): 0.6326 / 0.0078
[1/10][0/28]   Loss_D: 0.8028  Loss_G: 18.4099  D(x): 0.9846  D(G(z)): 0.4647 / 0.0000
[2/10][0/28]   Loss_D: 0.2333  Loss_G: 10.0259  D(x): 0.9677  D(G(z)): 0.1397 / 0.0001
[3/10][0/28]   Loss_D: 0.8398  Loss_G: 20.5674  D(x): 0.9399  D(G(z)): 0.4162 / 0.0000
[4/10][0/28]   Loss_D: 1.0563  Loss_G: 15.2354  D(x): 0.9563  D(G(z)): 0.4299 / 0.0000
[5/10][0/28]   Loss_D: 0.4594  Loss_G: 7.5027  D(x): 0.9790  D(G(z)): 0.3057 / 0.0015
[6/10][0/28]   Loss_D: 2.4790  Loss_G: 6.4899  D(x): 0.2710  D(G(z)): 0.0020 / 0.0136
[7/10][0/28]   Loss_D: 0.6013  Loss_G: 5.8628  D(x): 0.7983  D(G(z)): 0.1771 / 0.0058
[8/10][0/28]   Loss_D: 2.2355  Loss_G: 11.7234  D(x): 0.9178  D(G(z)): 0.7084 / 0.0001
[9/10][0/28]   Loss_D: 1.0205  Loss_G: 6.8938  D(x): 0.9690  D(G(z)): 0.4207 / 0.0054
```



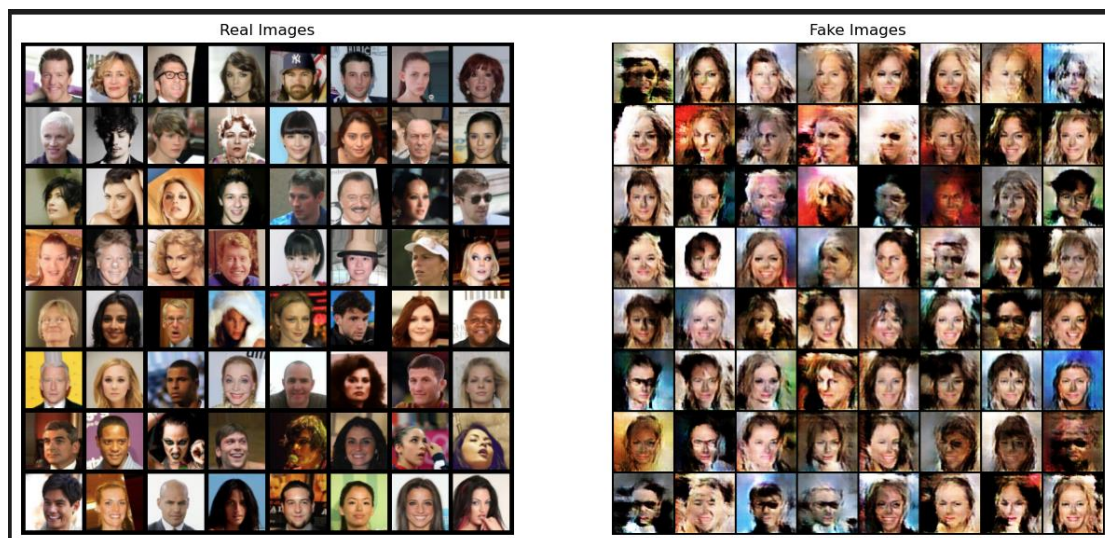

Generator 的 loss 隨著 epochs 數目慢慢下降。

Batch size = 128, Epochs = 100, Learning Rate = 0.0002


```

Output exceeds the size limit. Open the full output data in a text editor
Starting Training Loop...
[0/100][0/28] Loss_D: 2.0722 Loss_G: 5.2566 D(x): 0.4722 D(G(z)): 0.6326 / 0.0078
[1/100][0/28] Loss_D: 2.6923 Loss_G: 24.7116 D(x): 0.9891 D(G(z)): 0.8790 / 0.0000
[2/100][0/28] Loss_D: 0.2255 Loss_G: 25.3161 D(x): 0.8786 D(G(z)): 0.0000 / 0.0000
[3/100][0/28] Loss_D: 0.2066 Loss_G: 10.5084 D(x): 0.9348 D(G(z)): 0.0153 / 0.0001
[4/100][0/28] Loss_D: 0.4519 Loss_G: 7.3841 D(x): 0.8803 D(G(z)): 0.1769 / 0.0018
[5/100][0/28] Loss_D: 1.4655 Loss_G: 6.0187 D(x): 0.4774 D(G(z)): 0.0027 / 0.0108
[6/100][0/28] Loss_D: 0.2775 Loss_G: 5.7865 D(x): 0.8848 D(G(z)): 0.1001 / 0.0128
[7/100][0/28] Loss_D: 0.5584 Loss_G: 4.3664 D(x): 0.7380 D(G(z)): 0.0454 / 0.0315
[8/100][0/28] Loss_D: 2.2249 Loss_G: 13.1141 D(x): 0.9646 D(G(z)): 0.7776 / 0.0000
[9/100][0/28] Loss_D: 1.3330 Loss_G: 4.4361 D(x): 0.4548 D(G(z)): 0.0093 / 0.0566
[10/100][0/28] Loss_D: 1.1526 Loss_G: 7.2297 D(x): 0.8753 D(G(z)): 0.5036 / 0.0024
[11/100][0/28] Loss_D: 3.0837 Loss_G: 11.3218 D(x): 0.9888 D(G(z)): 0.8525 / 0.0001
[12/100][0/28] Loss_D: 0.9804 Loss_G: 3.7104 D(x): 0.5164 D(G(z)): 0.0291 / 0.1069
[13/100][0/28] Loss_D: 4.2826 Loss_G: 2.5567 D(x): 0.0880 D(G(z)): 0.0003 / 0.2559
[14/100][0/28] Loss_D: 1.0059 Loss_G: 5.7475 D(x): 0.9442 D(G(z)): 0.4671 / 0.0217
[15/100][0/28] Loss_D: 0.6190 Loss_G: 8.0287 D(x): 0.9251 D(G(z)): 0.3562 / 0.0010
[16/100][0/28] Loss_D: 0.3446 Loss_G: 6.0492 D(x): 0.9319 D(G(z)): 0.1930 / 0.0069
[17/100][0/28] Loss_D: 0.5325 Loss_G: 5.6548 D(x): 0.7726 D(G(z)): 0.1407 / 0.0102
[18/100][0/28] Loss_D: 0.5556 Loss_G: 4.1307 D(x): 0.8870 D(G(z)): 0.2521 / 0.0358
[19/100][0/28] Loss_D: 2.7596 Loss_G: 6.7652 D(x): 0.2153 D(G(z)): 0.0047 / 0.0271
[20/100][0/28] Loss_D: 1.1223 Loss_G: 8.7262 D(x): 0.9319 D(G(z)): 0.5685 / 0.0009
[21/100][0/28] Loss_D: 2.7603 Loss_G: 14.3044 D(x): 0.9642 D(G(z)): 0.8699 / 0.0002
[22/100][0/28] Loss_D: 0.8788 Loss_G: 7.7616 D(x): 0.9328 D(G(z)): 0.4173 / 0.0017
[23/100][0/28] Loss_D: 0.4654 Loss_G: 4.0891 D(x): 0.7560 D(G(z)): 0.0689 / 0.0569
...
[96/100][0/28] Loss_D: 0.5529 Loss_G: 5.4246 D(x): 0.9472 D(G(z)): 0.3262 / 0.0122
[97/100][0/28] Loss_D: 1.2658 Loss_G: 6.6875 D(x): 0.9548 D(G(z)): 0.5331 / 0.0029
[98/100][0/28] Loss_D: 0.2353 Loss_G: 4.0446 D(x): 0.9533 D(G(z)): 0.1548 / 0.0272
[99/100][0/28] Loss_D: 3.3328 Loss_G: 8.0434 D(x): 0.9970 D(G(z)): 0.8887 / 0.0026

```




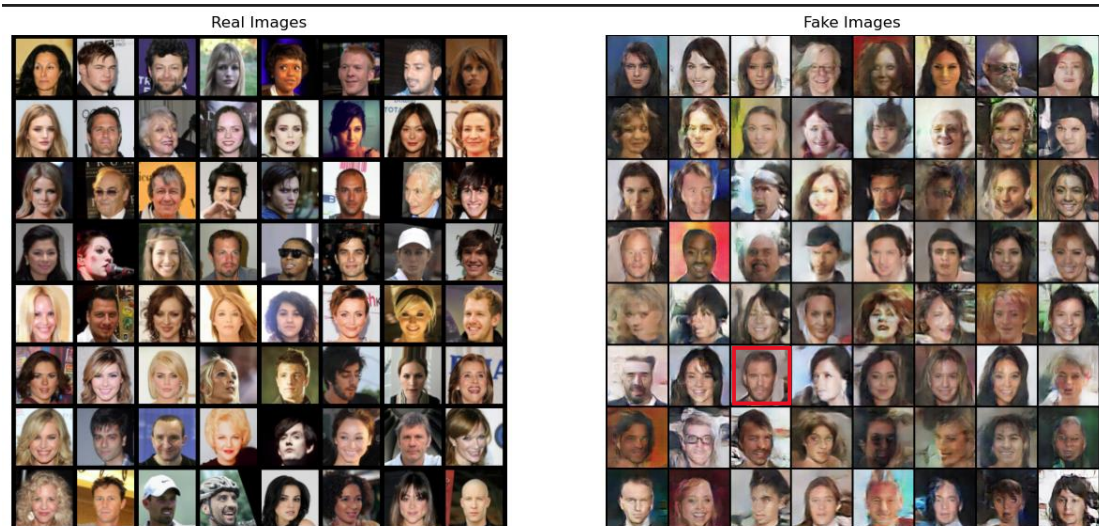
經過 100 次 epochs 後，可以明顯看出 generator 產生的 fake image 基本上已經看得出男女與一些特徵像頭髮顏色等.....，但是

五官細節還沒有到非常清楚。

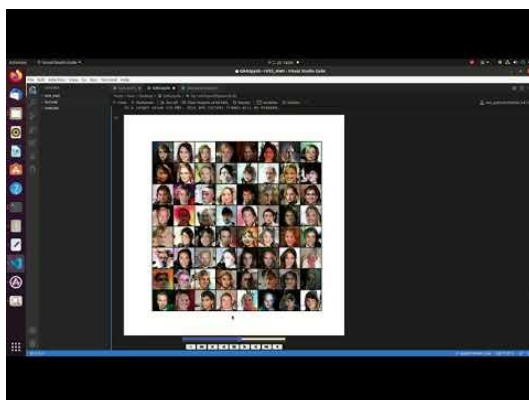
Batch size = 128, Epochs = 400, Learning Rate = 0.0002

```
Output exceeds the size limit. Open the full output data in a text editor
Starting Training Loop...
[0/400][0/28] Loss_D: 2.0722 Loss_G: 5.2564 D(x): 0.4722 D(G(z)): 0.6326 / 0.0078
[1/400][0/28] Loss_D: 0.0284 Loss_G: 7.5673 D(x): 0.9793 D(G(z)): 0.0030 / 0.0013
[2/400][0/28] Loss_D: 0.1725 Loss_G: 7.9746 D(x): 0.9496 D(G(z)): 0.0652 / 0.0017
[3/400][0/28] Loss_D: 0.3260 Loss_G: 16.2008 D(x): 0.8797 D(G(z)): 0.0001 / 0.0000
[4/400][0/28] Loss_D: 0.6842 Loss_G: 4.4837 D(x): 0.7180 D(G(z)): 0.0317 / 0.0245
[5/400][0/28] Loss_D: 0.4974 Loss_G: 1.7328 D(x): 0.7992 D(G(z)): 0.0914 / 0.2864
[6/400][0/28] Loss_D: 1.1975 Loss_G: 3.5966 D(x): 0.5057 D(G(z)): 0.0487 / 0.1144
[7/400][0/28] Loss_D: 3.0604 Loss_G: 13.2817 D(x): 0.9808 D(G(z)): 0.9160 / 0.0000
[8/400][0/28] Loss_D: 1.4099 Loss_G: 8.2310 D(x): 0.9281 D(G(z)): 0.5788 / 0.0009
[9/400][0/28] Loss_D: 0.5145 Loss_G: 5.8372 D(x): 0.8113 D(G(z)): 0.1676 / 0.0071
[10/400][0/28] Loss_D: 0.8261 Loss_G: 2.4659 D(x): 0.5842 D(G(z)): 0.0984 / 0.1165
[11/400][0/28] Loss_D: 1.8333 Loss_G: 11.6687 D(x): 0.9850 D(G(z)): 0.7382 / 0.0001
[12/400][0/28] Loss_D: 0.3852 Loss_G: 5.1070 D(x): 0.9043 D(G(z)): 0.2025 / 0.0228
[13/400][0/28] Loss_D: 3.4930 Loss_G: 6.2092 D(x): 0.9836 D(G(z)): 0.8749 / 0.0154
[14/400][0/28] Loss_D: 2.0229 Loss_G: 7.6743 D(x): 0.9622 D(G(z)): 0.7107 / 0.0108
[15/400][0/28] Loss_D: 0.7357 Loss_G: 3.5745 D(x): 0.6980 D(G(z)): 0.1941 / 0.0584
[16/400][0/28] Loss_D: 0.9289 Loss_G: 3.3970 D(x): 0.6167 D(G(z)): 0.1107 / 0.0949
[17/400][0/28] Loss_D: 0.8026 Loss_G: 5.8738 D(x): 0.9040 D(G(z)): 0.3404 / 0.0082
[18/400][0/28] Loss_D: 0.5578 Loss_G: 5.2229 D(x): 0.9066 D(G(z)): 0.3180 / 0.0142
[19/400][0/28] Loss_D: 0.4598 Loss_G: 4.0819 D(x): 0.7930 D(G(z)): 0.1340 / 0.0318
[20/400][0/28] Loss_D: 0.3599 Loss_G: 4.2459 D(x): 0.8418 D(G(z)): 0.1327 / 0.0257
[21/400][0/28] Loss_D: 2.3300 Loss_G: 14.1743 D(x): 0.9816 D(G(z)): 0.8243 / 0.0003
[22/400][0/28] Loss_D: 0.9257 Loss_G: 2.0608 D(x): 0.6276 D(G(z)): 0.1410 / 0.2109
[23/400][0/28] Loss_D: 0.3272 Loss_G: 2.8012 D(x): 0.8365 D(G(z)): 0.0953 / 0.1146
...
[396/400][0/28] Loss_D: 2.1372 Loss_G: 10.6567 D(x): 0.9991 D(G(z)): 0.7416 / 0.0002
[397/400][0/28] Loss_D: 1.3772 Loss_G: 9.3329 D(x): 0.9766 D(G(z)): 0.5427 / 0.0015
[398/400][0/28] Loss_D: 0.2459 Loss_G: 6.3186 D(x): 0.9856 D(G(z)): 0.1795 / 0.0035
[399/400][0/28] Loss_D: 0.1819 Loss_G: 5.2480 D(x): 0.9912 D(G(z)): 0.1452 / 0.0086
```





有幾張照片的五官已經非常清晰了，有無戴眼鏡或是有無鬍子，也都能看的出來，可是只有幾張照片而已，有些還是不能精確辨認出該照片的明顯特徵，而且 generator 的 loss 也似乎開始在悄悄上升，這邊有影片紀錄從第 1 個 epoch 到第 400 個的變化過程。



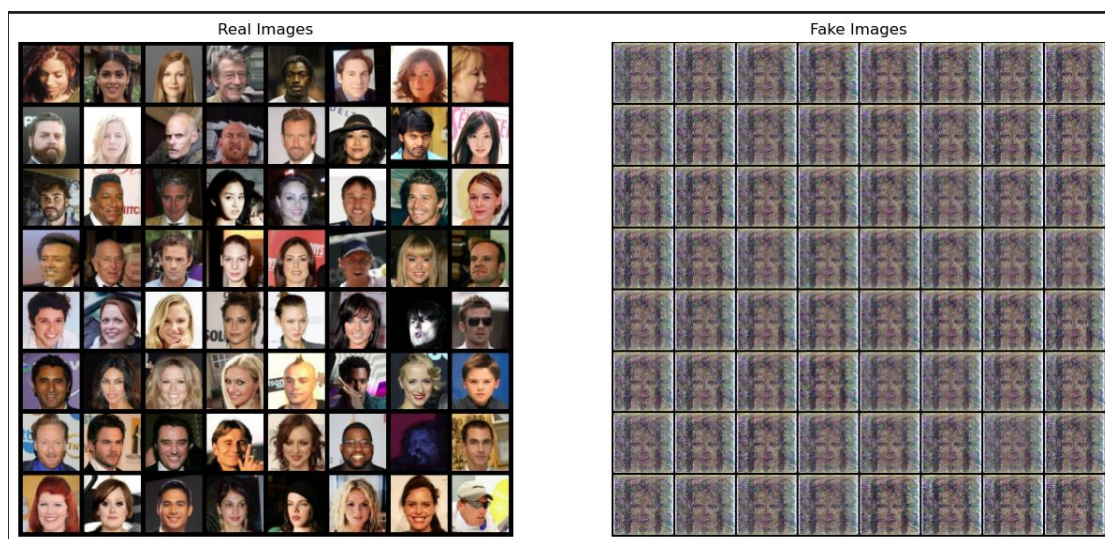

Batch size = 128, Epochs = 800, Learning Rate = 0.0002

Output exceeds the [size limit](#). Open the full output data [in a text editor](#)

Starting Training Loop...

```
[0/800][0/28] Loss_D: 2.0722 Loss_G: 5.2568 D(x): 0.4722 D(G(z)): 0.6326 / 0.0078
[1/800][0/28] Loss_D: 1.0333 Loss_G: 17.8370 D(x): 0.9936 D(G(z)): 0.5285 / 0.0000
[2/800][0/28] Loss_D: 0.0695 Loss_G: 9.4241 D(x): 0.9588 D(G(z)): 0.0053 / 0.0004
[3/800][0/28] Loss_D: 1.6803 Loss_G: 5.3873 D(x): 0.5005 D(G(z)): 0.0001 / 0.0234
[4/800][0/28] Loss_D: 1.2162 Loss_G: 13.3154 D(x): 0.9445 D(G(z)): 0.4768 / 0.0001
[5/800][0/28] Loss_D: 0.7829 Loss_G: 6.2318 D(x): 0.8069 D(G(z)): 0.2245 / 0.0082
[6/800][0/28] Loss_D: 0.5214 Loss_G: 4.7777 D(x): 0.7612 D(G(z)): 0.0959 / 0.0269
[7/800][0/28] Loss_D: 1.5593 Loss_G: 8.3328 D(x): 0.6082 D(G(z)): 0.3935 / 0.0005
[8/800][0/28] Loss_D: 1.2727 Loss_G: 8.2785 D(x): 0.9897 D(G(z)): 0.5761 / 0.0009
[9/800][0/28] Loss_D: 0.3628 Loss_G: 5.0411 D(x): 0.9499 D(G(z)): 0.2123 / 0.0214
[10/800][0/28] Loss_D: 1.2740 Loss_G: 3.5192 D(x): 0.8066 D(G(z)): 0.3345 / 0.0629
[11/800][0/28] Loss_D: 3.2811 Loss_G: 6.9405 D(x): 0.9922 D(G(z)): 0.9065 / 0.0086
[12/800][0/28] Loss_D: 0.6050 Loss_G: 6.2855 D(x): 0.8777 D(G(z)): 0.3242 / 0.0089
[13/800][0/28] Loss_D: 1.6046 Loss_G: 8.0623 D(x): 0.9760 D(G(z)): 0.6822 / 0.0016
[14/800][0/28] Loss_D: 1.9976 Loss_G: 9.4141 D(x): 0.9335 D(G(z)): 0.6832 / 0.0015
[15/800][0/28] Loss_D: 0.5581 Loss_G: 7.3354 D(x): 0.8963 D(G(z)): 0.3266 / 0.0017
[16/800][0/28] Loss_D: 0.5195 Loss_G: 5.7921 D(x): 0.8118 D(G(z)): 0.1468 / 0.0080
[17/800][0/28] Loss_D: 0.8322 Loss_G: 4.4292 D(x): 0.7979 D(G(z)): 0.2970 / 0.0300
[18/800][0/28] Loss_D: 0.8425 Loss_G: 6.0485 D(x): 0.9182 D(G(z)): 0.4190 / 0.0099
[19/800][0/28] Loss_D: 0.4095 Loss_G: 2.8562 D(x): 0.7618 D(G(z)): 0.0349 / 0.1054
[20/800][0/28] Loss_D: 0.5023 Loss_G: 7.4408 D(x): 0.9069 D(G(z)): 0.2435 / 0.0026
[21/800][0/28] Loss_D: 0.7940 Loss_G: 3.1732 D(x): 0.6871 D(G(z)): 0.2008 / 0.1033
[22/800][0/28] Loss_D: 0.4760 Loss_G: 5.6257 D(x): 0.8855 D(G(z)): 0.2304 / 0.0088
[23/800][0/28] Loss_D: 0.7638 Loss_G: 5.6807 D(x): 0.8761 D(G(z)): 0.3619 / 0.0110
...
```

```
[796/800][0/28] Loss_D: 0.0000 Loss_G: 86.3499 D(x): 1.0000 D(G(z)): 0.0000 / 0.0000
[797/800][0/28] Loss_D: 0.0000 Loss_G: 86.0833 D(x): 1.0000 D(G(z)): 0.0000 / 0.0000
[798/800][0/28] Loss_D: 0.0000 Loss_G: 86.6799 D(x): 1.0000 D(G(z)): 0.0000 / 0.0000
[799/800][0/28] Loss_D: 0.0000 Loss_G: 86.4647 D(x): 1.0000 D(G(z)): 0.0000 / 0.0000
```




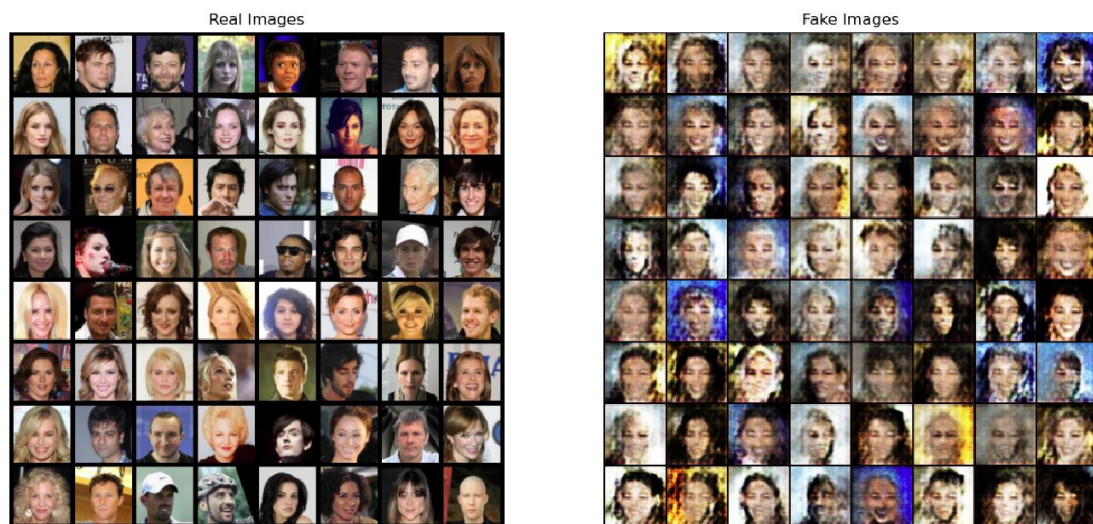
Generator 的 loss 直接爆長，以至於即使 epochs 數目很高也得不

出理想的 fake image，每張圖像甚至有一樣的結果，或許可能是 discriminator 訓練的太好了導致 generator 怎麼樣都會被抓包，因此無法一直進步。

Batch size = 2048, Epochs = 400, Learning Rate = 0.0002

```
Output exceeds the size limit. Open the full output data in a text editor
Starting Training Loop...
[0/400][0/2]   Loss_D: 1.9682  Loss_G: 4.7933  D(x): 0.4662  D(G(z)): 0.5937 / 0.0147
[1/400][0/2]   Loss_D: 0.6391  Loss_G: 6.7177  D(x): 0.9007  D(G(z)): 0.3391 / 0.0024
[2/400][0/2]   Loss_D: 0.8017  Loss_G: 7.1303  D(x): 0.8600  D(G(z)): 0.3611 / 0.0017
[3/400][0/2]   Loss_D: 0.5437  Loss_G: 7.5403  D(x): 0.8432  D(G(z)): 0.1593 / 0.0010
[4/400][0/2]   Loss_D: 0.3911  Loss_G: 8.2853  D(x): 0.8869  D(G(z)): 0.1275 / 0.0005
[5/400][0/2]   Loss_D: 0.2978  Loss_G: 8.5592  D(x): 0.8884  D(G(z)): 0.0550 / 0.0004
[6/400][0/2]   Loss_D: 0.2972  Loss_G: 9.5552  D(x): 0.8683  D(G(z)): 0.0148 / 0.0001
[7/400][0/2]   Loss_D: 0.2524  Loss_G: 9.5718  D(x): 0.9142  D(G(z)): 0.0605 / 0.0001
[8/400][0/2]   Loss_D: 0.2006  Loss_G: 10.4702 D(x): 0.9115  D(G(z)): 0.0092 / 0.0001
[9/400][0/2]   Loss_D: 0.2753  Loss_G: 11.4971 D(x): 0.9356  D(G(z)): 0.1192 / 0.0000
[10/400][0/2]  Loss_D: 0.4758  Loss_G: 16.4621 D(x): 0.9455  D(G(z)): 0.2878 / 0.0000
[11/400][0/2]  Loss_D: 0.1552  Loss_G: 8.5349  D(x): 0.9275  D(G(z)): 0.0026 / 0.0004
[12/400][0/2]  Loss_D: 0.3409  Loss_G: 22.3306 D(x): 0.8404  D(G(z)): 0.0000 / 0.0000
[13/400][0/2]  Loss_D: 0.1360  Loss_G: 7.7663  D(x): 0.9396  D(G(z)): 0.0012 / 0.0009
[14/400][0/2]  Loss_D: 0.3474  Loss_G: 23.0745 D(x): 0.8428  D(G(z)): 0.0000 / 0.0000
[15/400][0/2]  Loss_D: 0.1718  Loss_G: 9.1367  D(x): 0.9185  D(G(z)): 0.0003 / 0.0003
[16/400][0/2]  Loss_D: 0.2682  Loss_G: 19.0750 D(x): 0.8807  D(G(z)): 0.0000 / 0.0000
[17/400][0/2]  Loss_D: 0.1402  Loss_G: 5.8880  D(x): 0.9409  D(G(z)): 0.0115 / 0.0053
[18/400][0/2]  Loss_D: 0.7008  Loss_G: 26.2235 D(x): 0.7120  D(G(z)): 0.0000 / 0.0000
[19/400][0/2]  Loss_D: 0.1807  Loss_G: 14.9547 D(x): 0.9176  D(G(z)): 0.0000 / 0.0000
[20/400][0/2]  Loss_D: 2.8557  Loss_G: 20.9148 D(x): 0.9710  D(G(z)): 0.8997 / 0.0000
[21/400][0/2]  Loss_D: 0.4380  Loss_G: 20.3023 D(x): 0.8064  D(G(z)): 0.0000 / 0.0000
[22/400][0/2]  Loss_D: 0.1610  Loss_G: 4.9632  D(x): 0.9352  D(G(z)): 0.0132 / 0.0119
[23/400][0/2]  Loss_D: 0.5645  Loss_G: 24.3075 D(x): 0.7540  D(G(z)): 0.0000 / 0.0000
...
[396/400][0/2] Loss_D: 0.4758  Loss_G: 4.4719 D(x): 0.8962  D(G(z)): 0.2580 / 0.0220
[397/400][0/2] Loss_D: 0.4466  Loss_G: 4.3238 D(x): 0.8491  D(G(z)): 0.2093 / 0.0220
[398/400][0/2] Loss_D: 0.4434  Loss_G: 4.4555 D(x): 0.8833  D(G(z)): 0.2429 / 0.0194
[399/400][0/2] Loss_D: 0.5169  Loss_G: 4.8613 D(x): 0.8711  D(G(z)): 0.2786 / 0.0142
```






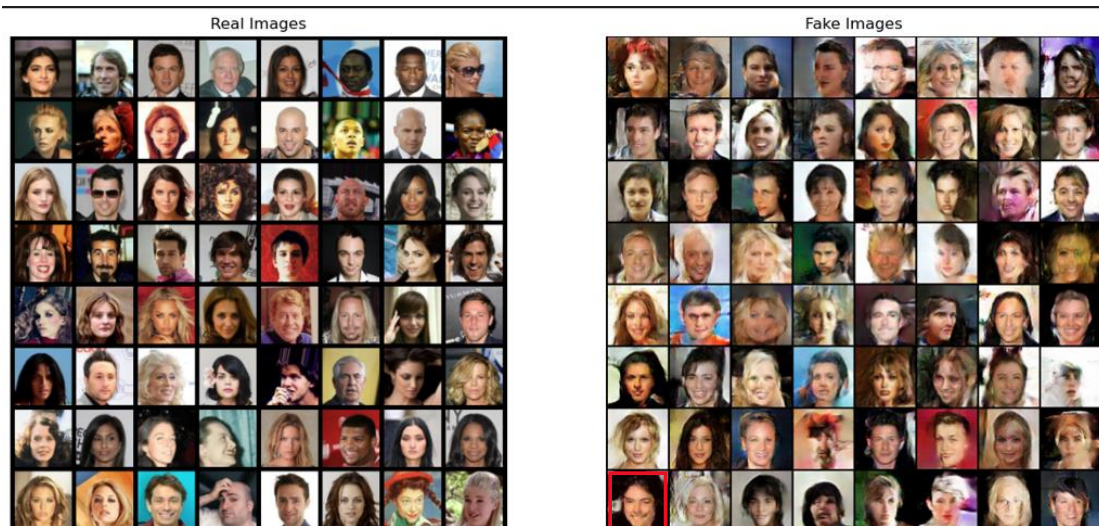
Batch size = 2048, Epochs = 2000, Learning Rate = 0.0002

```

Output exceeds the size limit. Open the full output data in a text editor
Starting Training Loop...
[0/2000][0/2] Loss_D: 1.9682 Loss_G: 4.7932 D(x): 0.4662 D(G(z)): 0.5937 / 0.0147
[1/2000][0/2] Loss_D: 0.6391 Loss_G: 6.7177 D(x): 0.9006 D(G(z)): 0.3392 / 0.0024
[2/2000][0/2] Loss_D: 0.8018 Loss_G: 7.1310 D(x): 0.8600 D(G(z)): 0.3611 / 0.0017
[3/2000][0/2] Loss_D: 0.5441 Loss_G: 7.5415 D(x): 0.8433 D(G(z)): 0.1598 / 0.0010
[4/2000][0/2] Loss_D: 0.3913 Loss_G: 8.2888 D(x): 0.8868 D(G(z)): 0.1276 / 0.0005
[5/2000][0/2] Loss_D: 0.2984 Loss_G: 8.5694 D(x): 0.8887 D(G(z)): 0.0560 / 0.0004
[6/2000][0/2] Loss_D: 0.2982 Loss_G: 9.5819 D(x): 0.8677 D(G(z)): 0.0144 / 0.0001
[7/2000][0/2] Loss_D: 0.2549 Loss_G: 9.6222 D(x): 0.9152 D(G(z)): 0.0644 / 0.0001
[8/2000][0/2] Loss_D: 0.1987 Loss_G: 10.2604 D(x): 0.9133 D(G(z)): 0.0111 / 0.0001
[9/2000][0/2] Loss_D: 0.2095 Loss_G: 9.8538 D(x): 0.9310 D(G(z)): 0.0565 / 0.0001
[10/2000][0/2] Loss_D: 0.1531 Loss_G: 10.5531 D(x): 0.9224 D(G(z)): 0.0060 / 0.0000
[11/2000][0/2] Loss_D: 0.1538 Loss_G: 10.0227 D(x): 0.9330 D(G(z)): 0.0142 / 0.0001
[12/2000][0/2] Loss_D: 0.1849 Loss_G: 13.3643 D(x): 0.9093 D(G(z)): 0.0007 / 0.0000
[13/2000][0/2] Loss_D: 1.8029 Loss_G: 23.5640 D(x): 0.9623 D(G(z)): 0.7729 / 0.0000
[14/2000][0/2] Loss_D: 0.3557 Loss_G: 21.3437 D(x): 0.8448 D(G(z)): 0.0000 / 0.0000
[15/2000][0/2] Loss_D: 0.3162 Loss_G: 12.2869 D(x): 0.9714 D(G(z)): 0.2051 / 0.0000
[16/2000][0/2] Loss_D: 0.2682 Loss_G: 12.6129 D(x): 0.9568 D(G(z)): 0.1471 / 0.0000
[17/2000][0/2] Loss_D: 0.4057 Loss_G: 16.2062 D(x): 0.9484 D(G(z)): 0.2390 / 0.0000
[18/2000][0/2] Loss_D: 0.1546 Loss_G: 7.6244 D(x): 0.9298 D(G(z)): 0.0058 / 0.0010
[19/2000][0/2] Loss_D: 0.6266 Loss_G: 26.8622 D(x): 0.7164 D(G(z)): 0.0000 / 0.0000
[20/2000][0/2] Loss_D: 0.1268 Loss_G: 14.0468 D(x): 0.9429 D(G(z)): 0.0000 / 0.0000
[21/2000][0/2] Loss_D: 3.7469 Loss_G: 25.0590 D(x): 0.9779 D(G(z)): 0.9587 / 0.0000
[22/2000][0/2] Loss_D: 0.4572 Loss_G: 26.3590 D(x): 0.8025 D(G(z)): 0.0000 / 0.0000
[23/2000][0/2] Loss_D: 0.0943 Loss_G: 10.4925 D(x): 0.9583 D(G(z)): 0.0001 / 0.0002
...
[1996/2000][0/2] Loss_D: 0.2356 Loss_G: 2.7955 D(x): 0.9050 D(G(z)): 0.1218 / 0.0801
[1997/2000][0/2] Loss_D: 0.2380 Loss_G: 2.8468 D(x): 0.9097 D(G(z)): 0.1283 / 0.0757
[1998/2000][0/2] Loss_D: 0.2408 Loss_G: 2.7624 D(x): 0.9019 D(G(z)): 0.1233 / 0.0830
[1999/2000][0/2] Loss_D: 0.2221 Loss_G: 2.8421 D(x): 0.9097 D(G(z)): 0.1152 / 0.0762

```






調高 batch size 後，fake image 看起來像是畫作，不過也蠻接近真實圖像的。

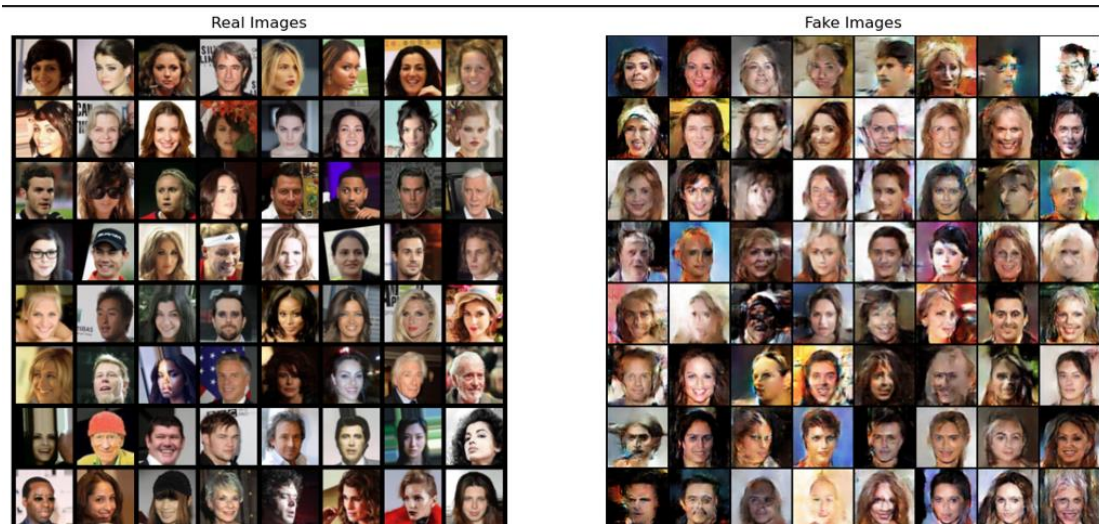
Batch size = 4096, Epochs = 5000, Learning Rate = 0.0001

```

Output exceeds the size limit. Open the full output data in a text editor
Starting Training Loop...
[0/5000][0/1] Loss_D: 1.9760 Loss_G: 2.3245 D(x): 0.4652 D(G(z)): 0.5930 / 0.1428
[1/5000][0/1] Loss_D: 1.3850 Loss_G: 2.9006 D(x): 0.8608 D(G(z)): 0.6347 / 0.0850
[2/5000][0/1] Loss_D: 0.9654 Loss_G: 3.5524 D(x): 0.8643 D(G(z)): 0.4847 / 0.0466
[3/5000][0/1] Loss_D: 0.6933 Loss_G: 3.9343 D(x): 0.8396 D(G(z)): 0.3304 / 0.0323
[4/5000][0/1] Loss_D: 0.5855 Loss_G: 4.0086 D(x): 0.8367 D(G(z)): 0.2576 / 0.0305
[5/5000][0/1] Loss_D: 0.5795 Loss_G: 4.0127 D(x): 0.8566 D(G(z)): 0.2722 / 0.0302
[6/5000][0/1] Loss_D: 0.6023 Loss_G: 4.2193 D(x): 0.8714 D(G(z)): 0.2968 / 0.0251
[7/5000][0/1] Loss_D: 0.5848 Loss_G: 4.4929 D(x): 0.8723 D(G(z)): 0.2834 / 0.0192
[8/5000][0/1] Loss_D: 0.5692 Loss_G: 4.7227 D(x): 0.8673 D(G(z)): 0.2638 / 0.0151
[9/5000][0/1] Loss_D: 0.5261 Loss_G: 4.9066 D(x): 0.8635 D(G(z)): 0.2272 / 0.0124
[10/5000][0/1] Loss_D: 0.4998 Loss_G: 5.0520 D(x): 0.8678 D(G(z)): 0.2104 / 0.0107
[11/5000][0/1] Loss_D: 0.4433 Loss_G: 5.1713 D(x): 0.8759 D(G(z)): 0.1789 / 0.0093
[12/5000][0/1] Loss_D: 0.4118 Loss_G: 5.2883 D(x): 0.8872 D(G(z)): 0.1698 / 0.0084
[13/5000][0/1] Loss_D: 0.3887 Loss_G: 5.4678 D(x): 0.8960 D(G(z)): 0.1637 / 0.0072
[14/5000][0/1] Loss_D: 0.3705 Loss_G: 5.6491 D(x): 0.9010 D(G(z)): 0.1563 / 0.0061
[15/5000][0/1] Loss_D: 0.3542 Loss_G: 5.7845 D(x): 0.9034 D(G(z)): 0.1479 / 0.0054
[16/5000][0/1] Loss_D: 0.3438 Loss_G: 5.9068 D(x): 0.9049 D(G(z)): 0.1418 / 0.0049
[17/5000][0/1] Loss_D: 0.3334 Loss_G: 6.0302 D(x): 0.9063 D(G(z)): 0.1364 / 0.0043
[18/5000][0/1] Loss_D: 0.3215 Loss_G: 6.1789 D(x): 0.9081 D(G(z)): 0.1292 / 0.0039
[19/5000][0/1] Loss_D: 0.3058 Loss_G: 6.2052 D(x): 0.9109 D(G(z)): 0.1227 / 0.0033
[20/5000][0/1] Loss_D: 0.2969 Loss_G: 6.4076 D(x): 0.9147 D(G(z)): 0.1220 / 0.0029
[21/5000][0/1] Loss_D: 0.2894 Loss_G: 6.5998 D(x): 0.9179 D(G(z)): 0.1211 / 0.0025
[22/5000][0/1] Loss_D: 0.2736 Loss_G: 6.7303 D(x): 0.9202 D(G(z)): 0.1126 / 0.0021
[23/5000][0/1] Loss_D: 0.2654 Loss_G: 6.8576 D(x): 0.9229 D(G(z)): 0.1099 / 0.0019
...
[4996/5000][0/1] Loss_D: 0.5228 Loss_G: 3.0659 D(x): 0.9052 D(G(z)): 0.3200 / 0.0641
[4997/5000][0/1] Loss_D: 0.4406 Loss_G: 2.0831 D(x): 0.7321 D(G(z)): 0.1014 / 0.1625
[4998/5000][0/1] Loss_D: 0.4049 Loss_G: 2.4865 D(x): 0.8804 D(G(z)): 0.2266 / 0.1076
[4999/5000][0/1] Loss_D: 0.3614 Loss_G: 2.5139 D(x): 0.8307 D(G(z)): 0.1504 / 0.1066

```





五、心得

我覺得訓練模型是一件很難的事，因為你永遠不知道怎樣可以
以得到最好的結果，它中間涉及的變因太多了不管是 optimizer 的選
擇或是 learning rate 的調整甚至是模型的組成架構。因此我以 loss 與
輸出的結果來調整 model 看看怎樣可以獲得完美的 fake image，據我
所知有一個叫做 optuna 的工具可以幫忙找出最適合的
hyperparameters，以後可以使用這個工具，把我得 GAN 模型 train 的
更完美，訓練模型的過程除了花時間之外，更重要的是適時調整模
型與不斷地觀察 loss 等，看是否有梯度消失或是梯度爆炸甚至是
overfitting 的問題，因此我認為把一個模型訓練得很好，其實不是一
件很容易的事。

根據我的實驗，我認為在 (Batch size = 128, Epochs = 400, Learning

Rate = 0.0002)和(Batch size = 2048, Epochs = 2000, Learning Rate = 0.0002)的時候生成出的 fake image 挺接近 real image 的。

這裡我提交的 ipynb 檔只有跑過 5 個 epochs 的紀錄，剩下的實驗是用實驗室的 server 來跑的，因為規格好非常多，可以節省非常多的時間。